Report: Image Segmentation Model

Name : Gudeme Nikhil Chand
Roll : 21f1003825

1. Dataset Description

For this image segmentation task, we chose the Lung CT Segmentation dataset from Kaggle. This dataset contains 2D lung CT scan images along with corresponding ground truth masks for lung regions. The goal of this dataset is to segment the lung region from the surrounding anatomy in CT scan images, which is useful for medical diagnosis and treatment planning.

- Dataset Size: The dataset consists of approximately 267 training images and masks.
- Image Dimensions: Each CT scan image is of size 512x512 pixels, saved in PNG format.
- Labels: The images are labeled with binary masks where the lung region is marked as 1 (foreground), and the background is labeled as 0.

The dataset provides high-resolution CT images, which makes it ideal for testing advanced segmentation models in medical imaging.

2. Model Used

For this project, we implemented a U-Net architecture, which is highly suited for biomedical image segmentation due to its ability to capture both global and local context.

- U-Net Architecture: The U-Net model is a fully convolutional network (FCN) with an encoder-decoder structure:
    - Encoder: Extracts feature maps using convolutional layers followed by max-pooling operations for downsampling.
    - Decoder: Upsamples the feature maps using transposed convolution layers and combines them with high-resolution features from the encoder via skip connections.

The U-Net model was chosen because of its strong performance in medical image segmentation tasks and its ability to accurately delineate boundaries of small objects (like lung regions).

- Input size: (128, 128, 3) for reduced computational complexity.
- Output size: (128, 128, 1), which represents a binary segmentation mask.

3. Evaluation Results

The model was trained for 20 epochs using the Adam optimizer with a learning rate of 0.001 and binary cross-entropy as the loss function. The following metrics were used to evaluate the model's performance:

- Intersection over Union (IoU): IoU measures the overlap between the predicted mask and the ground truth mask. It is calculated as the ratio of the intersection of the true and predicted masks to their union. Our model achieved an IoU score of 0.87, which indicates good segmentation accuracy for lung regions.
- Dice Coefficient: The Dice coefficient is another metric for evaluating segmentation tasks, which measures the similarity between the predicted and ground truth masks. Our U-Net model obtained a Dice Coefficient of 0.91, signifying a high overlap between the predicted and actual lung regions.

4. Visualization of Results

We visualized the model’s predictions by comparing the original CT scans with the segmented lung masks. Below are some observations:

- The U-Net model was able to effectively segment the lung regions in most cases, showing clear boundaries between the lungs and surrounding anatomy.
- The model struggled slightly with smaller features or images with low contrast, where the lung boundaries were less clear.

Example of segmented output:

- Left: Original CT scan
- Right: Segmented lung mask

5. Conclusion

The U-Net model provided reliable segmentation results on the lung CT dataset, achieving high accuracy with both IoU and Dice Coefficient metrics. The results indicate that the model is capable of segmenting lung regions with minimal error, making it suitable for medical imaging tasks. For further improvement, using larger input dimensions and increasing the model depth could enhance its performance, especially in detecting finer details.